

Donghyung Lee

DevOps Engineer | Resume & Career Description

E-mail: genius5711@gmail.com Blog(KR): somaz.tistory.com Blog(EN): somaz.blog GitHub: github.com/somaz94

Introduction

Hello, I am Donghyung Lee, a DevOps Engineer.

Targeting 99.9% infra availability, I codify every change as IaC and turn incident response into postmortem-driven assets to raise operational efficiency. My core strength is interrogating the purpose of each task and pursuing a better direction.

I design and build infrastructure across AWS and GCP cloud and an OpenStack-based on-premises PaaS. I standardize container operations with Kubernetes, codify infrastructure with Terraform and Ansible, and run GitLab CI / ArgoCD-centered deploy pipelines alongside monitoring and logging stacks. Repetitive work is automated in Python and Bash.

I also operate adjacent data-engineering and AI/ML ops through data-pipeline automation and an in-house GPU LLM service, and build in-house DevOps tools myself — git-bridge, a Slack APK bot, and more.

Across 3 companies I solely led AWS → GCP migration, AWS + on-premises hybrid, and on-premises PaaS builds — achieving 20% AWS and 50% GCP cost cuts, 50-67% faster deploys, and 90%+ manual-work automation. Most recently I resolved 14-second game-DB stalls, cut CI bootstrap time from 5min to 10s, and unified login via Keycloak central SSO — raising operational efficiency another notch.

I aim to keep contributing to infrastructure architecture advancement and organizational productivity.

Last updated: 2026.05

Strengths

Asking 'Why' First

- Habit of asking 'why' first — define the business problem before picking the tool

Systematic Documentation & Knowledge Sharing

- Document all work processes and share with team, contributing to organizational knowledge assets

Record-Based Continuous Improvement

- Postmortem-record every incident, infra change, and troubleshooting — preventing recurrence, driving improvement

Full-Stack Infrastructure Experience

- End-to-end infrastructure experience spanning cloud (AWS, GCP), on-premises, networking, CI/CD, and monitoring

Skills

Cloud	AWS, GCP, OpenStack
Container Orchestration	Kubernetes (EKS/GKE Multi-cluster Ops, On-premises HA with CNI/CSI Optimization), Cilium CNI (eBPF, NetworkPolicy zero-trust), NGINX Gateway Fabric / Gateway API
Infrastructure as Code	Terraform, Ansible, Helm, Helmfile
CI/CD	GitHub Actions, GitLab CI/CD, ArgoCD, Jenkins
Monitoring & Logging	PLG Stack, ELK Stack, EFK Stack, ECK Operator, Fluent Bit, Thanos
Security & IAM	Keycloak (Operator, OIDC IdP), Vaultwarden, GitLab SSO (OIDC)
Automation	Go, Python, Shell Scripting
Database	MySQL, PostgreSQL, Redis, MongoDB
OS & Servers	Ubuntu, Rocky Linux, CentOS, Debian, GPU Server

Core Competencies

Cloud Infrastructure Design & Operations

Multi/hybrid cloud architecture design and operations experience on AWS and GCP

- AWS EKS, GCP GKE production operations & cost optimization (20% reduction via hybrid transition, 50% via GCP migration, savings reinvested in new projects)
- IaC-based infrastructure automation and version control using Terraform (100% GCP resource codification)
- On-premises to cloud hybrid architecture design and migration experience
- Kubernetes storage performance optimization (control-plane HDD → SSD migration — etcd fsync 1760ms → 3.88ms = 454× faster; GitLab Runner build I/O isolation recovered game DB COMMIT latency from 14s back to ms)

CI/CD Pipeline & GitOps

Maximizing development productivity through GitOps-based deployment automation

- CI/CD pipeline design and enhancement using GitHub Actions, GitLab CI, ArgoCD
- 50-67% deployment time reduction with 3x weekly deployments shortening Time to Market, 95% rollback time decrease (30min → 1-2min)
- Jenkins-based Unity Android/iOS build automation and Slack-integrated deployment notifications
- Built a layered shared toolchain image architecture + auto version cycle standardizing 15 consumer repos — first-job install time 5min → ~10s (~95% reduction), with 4-stage cascade automated by just 2 manual clicks

Monitoring & Observability System Implementation

Establishing proactive incident detection and rapid response systems

- Integrated metric/log monitoring with PLG Stack (Prometheus, Loki, Grafana)
- Evolved from ELK → EFK Stack, then redesigned into ECK Operator-based Elastic Stack 9.0 CRD declarative management (automatic TLS rotation, automated rolling upgrades)
- Reduced MTTR and improved service availability via SLI/SLO-based alerting
- kube-prometheus-stack with custom alert rules, Grafana dashboards, and 4 DB exporters (MySQL, Redis, Elasticsearch, PostgreSQL) integrated
- Fluent Bit data integrity & HA hardening (SQLite offset DB + PVC, PodDisruptionBudget + preStop hook, 7 Prometheus alerts — zero log gap / duplication on Pod restart)

Automation & Efficiency

Operations automation and developer productivity tools using Python and Bash

- 90%+ manual work reduction through data collection, deployment, and document management automation
- Internal LLM service (Ollama + Open WebUI) ensuring 100% internal security guideline compliance and dev team productivity improvement
- Internal tool development: GitLab Webhook doc sync, Slack Bot APK deployment notifications
- Migrated shell → Python (stdlib only, 654 unittests) + wave-based auto-deploy across 25 components + auto-upgrade MR generation
- Established self-sufficient Kubernetes ops at DP (Delivery Partner) sites — designed and ran a 6-month training program (avg 10 attendees, 4h/session), standardized materials reused across other clients

Kubernetes Platform Modernization & Open-Source Contribution

Zero-downtime migration of network/logging stacks and releasing reusable public charts as open source

- Migrated multiple ingress-nginx instances to a single NGINX Gateway Fabric control-plane + per-class Gateway CRs with zero downtime (full cutover, wildcard TLS consolidated)
- Migrated the ECK Operator-based logging stack from a local chart to an OCI chart with zero downtime (CR/PVC names preserved, cluster_uuid unchanged, no data loss)
- Released 3 open-source Helm charts (nginx-gateway-cr, elasticsearch-eck, kibana-eck) and multiple composite GitHub Actions on ArtifactHub / GitHub Marketplace

Infrastructure Security & IAM

Centralized in-house authentication and secret governance via Keycloak Operator-based SSO and Vaultwarden password management

- Introduced central SSO with the Keycloak Operator + KeycloakRealmImport — broker GitLab so existing accounts keep logging in, integrated ArgoCD / Harbor OIDC (zero downtime, verification checks PASS), and established a single authorization flow from GitLab groups → Keycloak mappers → ArgoCD roles
- Deployed Vaultwarden on Kubernetes via Helm with GitLab SSO + automated backups (SQLite dumps, 30-day retention) + interactive restore script — centralized management for 70 users and 50 secrets
- git-bridge Retry API — zero webhook-secret exposure + SRE-friendly recovery (Bearer auth, repo-level direction override)

Data Engineering & Internal LLM Services

End-to-end ownership from GCP data pipelines to in-house LLM services

- Built a multi-source (MongoDB · CloudSQL · GA · Dune) → BigQuery → Google Sheets pipeline by combining BigQuery, Dataflow, Cloud Functions, and Cloud Scheduler — entire infrastructure 100% Terraform IaC
- Deployed an internal LLM service (Ollama + Open WebUI) on a Kubernetes-hosted on-premises GPU server, saving ~\$100/month by replacing external AI SaaS subscriptions and ensuring 100% internal security compliance
- 95% data-collection automation and 0% data gap rate by eliminating human error

Professional Experience (Detailed)

Phantom (Concrit Studio)

2024.10 ~ Present

DevOps Engineer

As the sole DevOps infrastructure engineer at a game development company, I design, build, and operate AWS-based cloud and on-premises servers. Development and test servers are on-premises, while QA and Production servers are on AWS, operating a hybrid architecture that reduced monthly infrastructure costs by 20% (\$60K → \$48K) and channeled the savings into new project investments. Currently providing secondary technical support and operations for a live-service game in collaboration with the external publisher, and building on-premises/AWS infrastructure to improve development productivity for new game projects.

Project 1: AWS-On-premises Hybrid Cloud Architecture

Contribution 100% (Solo design & implementation, in collaboration with the external publisher's ops team) 2024.10 ~ Present (In Operation)

Problem

While operating the entire infrastructure on AWS, high cloud costs of \$60,000/month necessitated cost optimization.

Approach

Workload-specific cost analysis revealed that dev/test environments had minimal traffic variation, making cloud auto-scaling unnecessary, while only QA/Production required scalability and stability. Designed a hybrid architecture transitioning dev/test to on-premises while keeping QA/Production on AWS.

The two environments are operationally independent without VPN (application traffic isolated); the image registry is also split by environment (dev = Harbor, QA/Production = ECR), with each environment building independently to match its architecture. Unified operational standards through Terraform and Ansible IaC integration (100% IaC coverage for new components, IAM excluded) across both environments, achieving consistent infrastructure management without increasing operational complexity.

Troubleshooting

- When configuring ALB Ingress on EKS, needed to route to different backends per game client version.
Resolved by combining custom header-based routing rules with Target Groups in ALB conditional routing
- Intermittent timeouts during ClusterIP service communication in Kubernetes IPVS mode.
Resolved through Cilium eBPF packet flow analysis and IPVS conntrack threshold tuning, and built a Cilium Network Policy-based zero-trust security model to enhance inter-pod traffic visibility and security
- During rolling deploys, the ALB target deregistration delay (lowered from the default 300s to 60s) didn't line up with Pod SIGTERM handling, causing some game-session connection drops. Solved via 2 changes — achieving zero-downtime deploys with 0 connection drops:
 - (1) Added a 60s sleep in the preStop hook (failing the readiness probe so ALB marks the target unhealthy and stops sending new requests)
 - (2) Set terminationGracePeriodSeconds to 90s (ALB drain 60s + in-flight handling 30s)

Results

- Workload analysis-based infra optimization reducing monthly costs by 20% (\$60,000 → \$48,000), reinvesting ~\$144,000 annual savings into new project infrastructure for resource reinvestment cycle
- Offloaded dev servers to on-premises, cutting ~\$1,000/month in AWS resource costs and optimizing resource usage
- Ensured operational stability and fault isolation through environment role separation

Tech Stack

AWS (EKS, EC2, ALB, Route53, ACM, EFS, Aurora MySQL, ElastiCache, CloudFront), Kubernetes, Terraform, Helm

Project 2: On-premises CI/CD Pipeline Construction & Enhancement

Contribution 100% (Full infra & pipeline ownership) 2024.12 ~ Present (In Operation)

Problem

Developers manually built and deployed to servers, taking an average of 30 minutes per deployment with frequent failures due to human errors.

Approach

Built a GitOps-based automated deployment environment combining GitLab CI and ArgoCD.

Eliminated Docker-in-Docker dependency using Kaniko for Docker image builds, and configured multi-environment (alpha/review/dev/staging) deployments in a single pipeline.

Troubleshooting

- ArgoCD Application's `syncPolicy.automated.selfHeal=true` reverted operators' temporary hot-fixes instantly, making incident debugging hard.
Split policy by environment (prod: `selfHeal=false` + manual sync windows; dev/staging: `selfHeal=true`) and integrated Slack alerts for every sync event to gain both change traceability and stability
- Package installation failure on CI Runner due to GPG key expiration after GitLab upgrade. Standardized the apt-key renewal procedure into a runbook for fast recovery on recurrence
- Build time spikes due to cache invalidation during Kaniko builds. Resolved with `-cache=true --cache-ttl=24h --snapshot-mode=redo` option combination

Results

- 67% deployment time reduction (30min → 10min), 90% manual work automation
- Deployment automation increased weekly deployments 3x+ (1-2 → 5-7), shortening Time to Market for new features, while deployment resource efficiency enabled the dev team to focus on core feature development

Tech Stack

GitLab CI/CD, ArgoCD, Jenkins, Kaniko, Kubernetes, Docker, Harbor

Project 3: Centralized Logging System (ELK → EFK → ECK Operator 3-stage Enhancement)

Contribution 100% (System design & implementation) 2024.11 ~ Present (ELK → EFK → ECK 3 Phase, ~18 months total)

Problem

Server and container logs were distributed across the on-premises environment, making root cause analysis time-consuming during incidents.

Troubleshooting

- Logstash parsing failures due to null bytes in game server logs. Resolved by pre-removing with mutate filter gsub
- Elasticsearch field mapping conflicts. Resolved with explicit index template mappings and type conversion filters
- Index read-only due to disk watermark exceeded. Stabilized with ILM policy auto-deleting indices older than 7 days
- [EFK] Fluent Bit NFS WAL write failure. Resolved with custom PVC template patch + upgrade script auto-preserve
- [EFK] Log loss due to Fluentd output buffer overflow. Resolved by tuning `chunk_limit_size` + `flush_interval` and setting `overflow_action=block`
- [ECK] Kibana 9 secure-cookie enforcement blocked HTTP login. Worked around with `tls.disabled` + `publicBaseUrl http://` + `secureCookies:false`
- [ECK] No Fluentd ES9 variant → ensured ES 9.0 compatibility with ES8 variant image via `fluent-plugin-elasticsearch 5.4.4` and `suppress_type_name:true`, verified end-to-end via active marker
- [ECK] 3.3.x `stackconfigpolicy-controller` repeatedly GET-ed resources outside managedNamespaces every 17 minutes, emitting reconcile errors.
Worked around by setting `managedNamespaces:[]` (cluster-wide watch); RBAC was already cluster-scoped so no permission delta

Approach

Phase 1 (ELK Setup): Built a centralized logging system with Elasticsearch + Logstash + Kibana + Filebeat. Filebeat collected node/container logs, Logstash handled JSON parsing and field mapping, Elasticsearch stored them, and Kibana dashboards provided real-time monitoring and search.

Phase 2 (EFK Migration): Migrated to a Fluent Bit + Fluentd EFK Stack to relieve Logstash's JVM memory burden (1GB+) and Filebeat's Helm-ecosystem mismatch. Lightweight C-based Fluent Bit ran as a DaemonSet with automatic Kubernetes metadata tagging, while Fluentd handled multi-output processing/filtering before forwarding to Elasticsearch. A chart-upgrade automation script preserved custom PVC templates across upstream chart auto-upgrades.

Phase 3 (ECK Operator + Stack 9.0): With Elastic's Helm Chart updates halted and CRD-based declarative management becoming mainstream, performed an 8.5.1 → 9.0.0 major upgrade and deployed ECK Operator 3.3.2 in the elastic-system namespace. Redefined ES/Kibana as Custom Resources (CR wrapper local chart), executed a zero-downtime cutover after parallel monitoring → logging ns deployment. Established automatic TLS rotation (1y / 30d), Helm-templated elastic account password, and rolling upgrades with a single-line CR spec.version change.

Results

- 90% log search time reduction, improved MTTR maintaining stable availability
- Removed Logstash's JVM memory burden (1GB+) via EFK migration (Fluent Bit + Fluentd architecture)
- 1GB daily log real-time processing, unified Helm management with automated upgrades

- ECK migration removed manual ops burden for TLS, passwords, StatefulSet upgrades; Stack 9.0 major upgrade closed a 2-year update gap
- Operator secureMode + ServiceMonitor auto-registers controller-runtime metrics into kube-prometheus-stack

Tech Stack

Elasticsearch, Fluent Bit, Fluentd, Logstash, Kibana, Filebeat, ECK Operator, CRD, Custom Resource, Kubernetes, Helm

Project 4: Developer Productivity Tools (APK Deploy Bot · Doc Automation · LLM Service · Git Mirroring · Static File Server)

Contribution 100% 2025.02 ~ Present

4-1. Slack APK Distribution Automation Bot (4 weeks, 2025.02)

Auto-notification to Slack with QR code on Jenkins build completion. 90% QA team delivery time reduction

Tech Stack

Python, Slack Bot API, Jenkins, Kubernetes, Docker

4-2. GitLab-Google Drive Document Automation System (2 weeks, 2025.06)

Auto-sync on push via GitLab Webhook + Google Drive API. 80% document management time reduction

Tech Stack

GitLab Webhook, Google Drive API, Python, Bash Script

4-3. Internal LLM Service Deployment (4 weeks, 2025.11)

Deployed Ollama + Open WebUI on Kubernetes on on-premises GPU server. Serves 15 developers with avg 100 queries/day, with an authored user guide and personally driven onboarding cementing adoption.

Saved ~\$100/month by replacing external AI SaaS subscriptions. 100% internal security guideline compliance and eliminated data leakage risks

Tech Stack

Ollama, Open WebUI, Kubernetes, NFS, GPU Server

4-4. Git Repository Mirroring Tool + Retry API Follow-up (4-week initial build, 2026.02 ~ 2026.06)

Developed a Go-based bidirectional Git mirroring tool (git-bridge). Ensured code quality with unit tests + a GitHub Actions CI/CD pipeline, fully automating sync across CodeCommit/GitLab/GitHub for 10 repos with avg 30 sync events/day.

Phase 2 (2026.05) added a `POST /retry/mirror` endpoint with Bearer-token auth (`RETRY_API_TOKEN`), separated from the webhook secret, constant-time compared) — a safe manual-retry path that no longer exposes the webhook secret. Direction supports `auto / source-to-target / target-to-source`, with per-repo `retry_direction` overrides and per-repo Slack webhook env vars for split notifications. Added 11 new tests, passing the CI test gate

Phase 3 (2026.06) root-caused a non-fast-forward regression where multiple uncoordinated force-pushes rolled a shared branch back to a stale snapshot, and designed `ref_overrides` to pin the sync direction per ref pattern. The repo stays bidirectional while specific branches are pinned one-way, preventing reverse-direction echoes from rolling a branch back; scoped pushes to the triggered ref (removed `--all`, applied only to override repos so existing repos behave exactly as before), added a direction restriction on delete events, and a fail-open fallback. Verified 6/6 runtime checks on a test-only repo (production repo isolated and unaffected), landing 4 commits (feat / fix / docs / chore) on main

Tech Stack

Go, AWS SQS, GitLab Webhook, GitHub Webhook, HTTP Bearer Auth, Slack Webhook, Kubernetes, Docker

4-5. Static File Server - In-house Development (Open Source, 2026.03 ~ Present)

Developed a lightweight Go-based static file server (static-file-server) in-house and operated a Helm repository on GitHub Pages. Features: dark-mode directory listing UI, file preview (image/video/PDF/text), search with URL-hash sharing, multi-select + ZIP batch download, Gzip, Prometheus metrics, JSON logging, automatic APK Content-Type.

Deployed to Kubernetes via an in-house Helm chart (NFS storage), fully replacing the legacy halverneus/static-file-server. Handles ~30 downloads/day and 5 APK builds/week (daily build). Released as open source under Apache-2.0.

Tech Stack

Go, Helm Chart, GitHub Pages, Docker Hub, Kubernetes, NFS, Prometheus

Project 5: Kubernetes Cluster Operations Enhancement

Contribution 100% (Solo design — monitoring · upgrade automation · Helm management framework · NGF migration · 3 OSS Helm charts · storage performance optimization · Shell→Python migration; captured as 4 automation scripts establishing a standard procedure → so the team can reuse) 2025.12 ~ Present

To advance cluster visibility, operational stability, network modernization, storage performance, and automation-code quality, overhauled monitoring, automated K8s upgrades and Helm management, and migrated the network controller onto the Gateway API. Released the artifacts from the NGF migration and ECK Operator adoption as OSS Helm charts, resolved storage bottlenecks via control-plane / DB-node SSD migration and build I/O isolation, and modernized the infra deploy/upgrade pipeline with wave-based automation + a Shell → Python migration.

5-1. Monitoring & Alerting Enhancement (2025.12 ~ Present)

Designed custom alert rules and Grafana dashboards based on kube-prometheus-stack and integrated MySQL/Redis/Elasticsearch/PostgreSQL exporters. Deployed node-exporter on physical servers and VMs via Ansible, and collected Cilium CNI agent/operator metrics via `kubernetes_sd_configs`. Configured severity-based Slack routing with inhibit rules in Alertmanager to improve alert quality. (Chose kube-prometheus-stack over SaaS Observability — Datadog / New Relic — because on-premises metric sovereignty, exporter freedom, and zero runtime cost mattered more.)

Tech Stack

Prometheus, Grafana, Alertmanager, Cilium, node-exporter, Ansible, Slack API

5-2. K8s Upgrade Automation & Helm Chart Management Framework (2025.12 ~ Present)

Kubespary upgrade automation script for pre-flight, backup, compatibility checks + multi-step health check script reducing verification time by 90%.

Upgrade framework across the Helm charts + 4 canonical templates synchronizer (check for CI drift detection, apply for batch propagation) maintaining script body consistency.

Automated Kubernetes certificate renewal script + crontab (across nodes, every 6 months at 3 AM) proactively preventing API Server outages from certificate expiration

(Kept Kubespary — kOps is AWS-only, Cluster API's on-premises bootstrap is heavyweight; Kubespary fits the Ansible-familiar ops, so only added the automation layer.)

Tech Stack

Kubespary, Ansible, Helm, etcd, Shell Script, GitLab CI

5-3. Ingress-nginx → NGINX Gateway Fabric (NGF) Migration (2026.04)

Migrated multiple ingress-nginx controller instances (default + public-a~j, each with a pinned MetalLB LoadBalancer IP) to a single NGF 2.x control-plane + per-class Gateway CRs, preserving MetalLB IPs so DNS/firewall stayed unchanged via parallel deployment and gradual cutover (30-60s actual downtime per class).

Executed Phase 0~7+ (MetalLB pool expansion → NGF install → observability apps → HTTP-only / ApplicationSet apps → HTTPS-enforced apps (Harbor, Vaultwarden) → IP-swap cutover → Ingress cleanup). Mapped ingress-nginx annotations 1:1 to Gateway API + NGF CRDs (HTTPRoute filters, ClientSettingsPolicy, ProxySettingsPolicy, RateLimitPolicy, BackendTLSPolicy) and unified the HTTPRoute values schema across the charts for batch migration of ApplicationSet apps.

Consolidated per-app self-signed TLS Secrets into a single wildcard TLS Secret (10-year validity), cutting manual renewal overhead by 50%.

Tech Stack

NGINX Gateway Fabric, Gateway API, HTTPRoute, ClientSettingsPolicy, ProxySettingsPolicy, RateLimitPolicy, BackendTLSPolicy, Helmfile, MetalLB, ArgoCD

5-4. Three Open-Source Helm Charts — `nginx-gateway-cr` · `elasticsearch-eck` · `kibana-eck` (2026.04 ~ 2026.05)

nginx-gateway-cr: refined the local CR chart extracted from the NGF migration (5-3) into a reusable chart that deploys five tenant-level resources — Gateway / NginxProxy / ReferenceGrant / ServiceMonitor / PodMonitor. Includes a multi-Gateway-friendly schema (`gateways[]` array with per-gateway overrides) and PodMonitor templates recommended for NGF 2.x (controller + dataplane).

elasticsearch-eck / kibana-eck: elasticsearch-eck (11 ES templates) + kibana-eck (10 Kibana templates) deploy the CRs, Secrets, HTTPRoute, Ingress, BackendTLSPolicy, ServiceMonitor, and NetworkPolicy from a single set of values. Added `values.schema.json` (draft-07) validation, Minimal/HA presets, and per-environment storageClass mapping in the README.

Unified distribution: all three charts ship simultaneously across OCI (ghcr.io/somaz94/charts) + a gh-pages Helm repo + ArtifactHub (networking / monitoring-logging, Apache-2.0). Added a new canonical version-tracking template.

Zero-downtime cutover: when migrating the mgmt cluster from the local chart to the OCI chart, every CR / PVC / Secret name was preserved — cluster_uuid stayed unchanged, with zero Elasticsearch pod restarts and a single Kibana rolling restart, i.e. no data loss. HA rolling scenarios passed 3/3 on a local kind cluster; end-to-end install on the real cluster reached Ready in 71s for Elasticsearch and 57s for Kibana.

Tech Stack

Helm Chart, Gateway API, NGINX Gateway Fabric, ECK Operator, Elasticsearch, Kibana, Prometheus Operator, OCI Registry, ArtifactHub, Kubernetes

5-5. Storage Performance Optimization — etcd / DB SSD Migration + Build I/O Isolation (2 weeks, 2026.04 ~ 2026.05)

The control-plane VM and the game-DB worker VM both ran on HDDs, contending with GitLab Runner buildx for I/O — etcd WAL fsync p99 ballooned to ~2s, apiserver pods GET p99 to ~7s, and game DB COMMIT latency to 9~14s, blocking play. Root cause pinned on the game-DB worker VM's saturated disk.

Step 1: tuned MySQL with 6 parameters (e.g. `innodb_flush_log_at_trx_commit=2`) for immediate relief.

Step 2 (control-plane VM SSD migration): sparse-copied the KVM qcow2 (11GB, 6m50s) + virsh XML edit, completing the cutover with 8.5 minutes of downtime (vs. a 15-minute estimate).

Step 3 (game-DB worker VM DB/Redis SSD migration): lossless migration of 35GB across game-db (MySQL) + 3-env valkey-redis via stop → backup → SSD migration → restore → local-path PV rebind.

Step 4 (build-only VM): provisioned a 4 vCPU / 16GB / 150GB SSD VM on a separate physical server via cloud-init, joined it with kubespray, and applied `node-role=build` label + `dedicated=build:NoSchedule` taint to isolate buildx. Captured the work as 4 automation scripts.

Postmortem & Recurrence Prevention: promoted the MAC sexagesimal pitfall into a MAC-match verification gate, promoted the etcd leader-election 5xx finding into an explicit 'etcd quorum recovery check' migration-runbook gate, and packaged the 4 automation scripts + config files as a standard procedure asset for future worker-node additions.

Results: etcd WAL fsync 1760ms → 3.88ms (454×), etcd backend commit 1810ms → 3.77ms (480×), apiserver pods GET p99 6520ms → 23ms (283×), game DB COMMIT 9~14s → ms range, build-isolation VM validated 11 PASS / 0 FAIL

Tech Stack

KVM / libvirt, qcow2, virsh, etcd, Prometheus / PromQL, MySQL, Redis (valkey), kubespray, cloud-init, GitLab Runner

5-6. Infra Deploy/Upgrade Automation + Shell → Python Migration (2026.03 ~ Present)

The internal Kubernetes infra monorepo's 25 components were managed by helmfile + shell, hitting limits on multi-env branching and complex YAML patching. **Deploy automation:** structured the CI pipeline into 6 waves (bootstrap → core → security → db-redis → observability → apps) — MR diff detection → helmfile diff/apply on changed components only, with a manual gate making the wave-based deploy a standard. Added 10 new CI scripts (helmfile-changed-dirs, diff-component, apply-component, auto-upgrade, render-mr-body, ...).

Upgrade automation: CI detects new versions across all 25 component charts and opens an MR with auto-rendered body (changed chart, values diff, breaking-change notes).

Shell → Python migration: rewrote 8 canonical templates + the sync and backup scripts + 6 CI scripts in Python (stdlib only). Built up 25 component-level upgrade modules + 654 unittests. Adopting the helmfile-tools image (from Project 7) brought **CI bootstrap time from 5min to ~10s (~95% reduction)**, with local (macOS venv) and CI container running the same code.

Tech Stack

Python 3.10+ (stdlib only), unittest, GitLab CI/CD, Helm / Helmfile, Git automation, yq

Project 6: Infrastructure Security System

Contribution 100% 2026.04 ~ Present

6-1. Vaultwarden Password Management System (1 week, 2026.04 ~ Present)

Deployed Vaultwarden on K8s via Helm, GitLab SSO integration, org/collection access control. Automated daily backup (30-day retention) + interactive restore.

Migrated 70 users and 50 secrets into centralized management

Tech Stack

Vaultwarden, Helm, OpenID Connect, GitLab SSO, Kubernetes

6-2. Keycloak Operator-based Central SSO Introduction (2026.04 ~ Present)

Replaced the layered setup where ArgoCD / Harbor / Vaultwarden each used GitLab directly as an OIDC IdP with a central broker (Keycloak). Introduced the Keycloak Operator + KeycloakCR + KeycloakRealmImport with a PostgreSQL backend, brokering GitLab as Identity Provider so existing accounts keep logging in. Codified Realm / Client / IdP / Group / Mapper via a `kcadm.sh` bootstrap script and migrated ArgoCD / Harbor's OIDC config from direct-GitLab to via-Keycloak (Vaultwarden held off). GitLab groups (`server` , `global-admin`) are brokered through Keycloak group mappers and wired straight into ArgoCD's role mappings, consolidating the authorization SSOT onto Keycloak.

Results: ArgoCD / Harbor OIDC integrations completed with zero downtime and verification checks PASS; captured five pitfalls hit during adoption (IdP providerId / scope / mapper config / Operator idempotency / Harbor user mapping) into the plan doc to prevent recurrence; laid the groundwork for future LDAP federation, MFA, and central session policy

Tech Stack

Keycloak Operator, KeycloakCR / KeycloakRealmImport, PostgreSQL, OIDC / OAuth2, kcadm.sh, ArgoCD, Harbor, Kubernetes Gateway API

Project 7: Shared Toolchain Image Layered Architecture + Auto Version Cycle

Contribution 100% (Architecture design · CI standardization · automation pipeline) 2025.09 ~ 2026.05

Problem

Across the 15 repos in the GitLab CI standardization scope (1 template provider + 14 consumers), each job re-installed helm / helmfile / kubectl / crane / aws-cli on top of an Alpine base, making first-job bootstrap take an average of ~5 minutes. Tool versions were scattered across the 14 consumer repos, making coordinated bumps effectively impossible.

Reliance on `:latest` tags also undermined reproducibility.

Approach

Layer 1 (Mirror): copied external registry images into the internal Harbor via `crane copy` , preserving multi-arch manifests.

Layer 2 (Custom): built 7 fat images on top of Layer 1 (helmfile-tools, node-build, go-build, rclone-backup, aws-cli-tools, kaniko-build, terraform-tools).

Tier 1~2 SSOT: introduced a shared `.toolchain_build_custom` template + central `TOOLCHAIN_*_TAG` variables in gitlab-ci-templates, indirecting consumer repos away from tracking image tags themselves.

Phase 4 auto cycle: a cron job drives the following 4 steps so the full cascade completes with just two manual clicks:

- (1) Upstream version detection
- (2) Layer 1 mirror auto-MR
- (3) Layer 2 rebuild + Tier 2 sync auto-MR
- (4) Auto-propagation to 14 consumer repos

Results

- **First-job bootstrap time cut from 5min to ~10s (~95% reduction)**, with 100% tool-version consistency across 14 consumer repos (15 repos total including the template provider)
- Compressed 12+ scattered version points into a 4-tier governance flow (ARG → central var → consumer indirection → lint drift check)
- Phase 4 cron now lands new versions of 5+ tools with just 2 manual clicks, effectively eliminating version-tracking overhead; the minor-guard policy prevents compatibility breakage in advance

Tech Stack

GitLab CI/CD, Docker Buildx (multi-arch), crane, Kaniko, Harbor, Helm / Helmfile, Bash / yq, Slack Webhook

Project 8: Fluent Bit Log Collector Data Integrity & HA Hardening

Contribution 100% (Tier 1~3 sole design & rollout — handed off to operators via 7 PrometheusRule annotations + Slack alert routing) 2026.03 ~ 2026.05

Problem

Operating the post-Phase-3 ECK stack revealed that whenever a Fluent Bit Pod restarted, its in-memory tail offsets were lost — risking either re-indexing the same logs or losing them entirely.

With only a single replica and no PDB / preStop hook, node drain caused brief collection gaps, and there were no alerts on the collector itself, leaving room for silent failure.

Solution

Tier 1 (data integrity): enabled a SQLite offset DB (`db /var/log/flb-storage/tail.db`) on the tail input and mounted a dedicated 2Gi PVC, so offsets survive Pod restarts. Verified on helmfile revision 4 — zero log re-index and zero duplicates.

Tier 2 (availability): added a PodDisruptionBudget (`minAvailable=1`), a preStop hook (`sleep + flush`), and liveness/readiness probes (`/api/v1/health`) to minimize collection gaps during node drain and rolling updates.

Tier 3 (monitoring): added 7 PrometheusRule alerts (input/output failure rates, retry buildup, fluentd queue pressure, fluent-bit Pod down, SQLite DB write failures, ...).

Tier 4 (HA option doc): documented three forward-looking HA options — input-path partitioning / sidecar collector / Vector multi-collector — with trade-off comparison.

Results

- Achieved zero log re-indexing and zero duplicates on Pod restart, guaranteeing log data integrity
- Minimized collection gaps during node drain / rolling update via PDB + preStop hook, contributing directly to MTTR
- Eliminated silent failures on the collector layer with 7 PrometheusRule alerts, establishing self-monitoring on the observability stack itself

Tech Stack

Fluent Bit (tail / SQLite offset), Kubernetes PVC / StorageClass, PodDisruptionBudget, Prometheus Operator (PrometheusRule), Helmfile, Fluentd (buffer tuning)

NerdyStar

2023.03 ~ 2024.07

DevOps Engineer

Worked as a DevOps Engineer at a game and blockchain startup, leading the large-scale cloud migration project from AWS to GCP.

Project 1: Large-scale AWS → GCP Cloud Migration & CI/CD Reconstruction

Contribution: Infra design/migration lead 70% (remaining 30% = server-team application-side migration collab), CI/CD pipeline 100% 10 months (2023.03 ~ 2023.12)

Problem

Continuously rising AWS costs (\$10,000+/month, high NAT Gateway and data transfer costs), a GCP Startup Program opportunity, and the need for GCP global load balancing for the multi-region game service

Approach

Adopted Shared VPC (Host / Service Project separation) and Workload Identity Federation so GitHub Actions could authenticate to GCP via short-lived tokens without service-account keys. Migrated DNS via subdomain delegation (Hosting.kr → AWS Route53 → GCP Cloud DNS) for pre-validation, then performed the final NS swap during a single scheduled maintenance window with minimal downtime.

Worked around Filestore's 2.5TB SSD minimum cost by self-hosting an NFS server on Compute Engine (pd-balanced 1TB). Configured Cloud Armor with region blocking + IP allow-list WAF policy, and split the Cloud CDN URL Map into HTTP / HTTPS variants to enforce HTTP → HTTPS redirect alongside static-asset caching.

Troubleshooting

- During AWS ALB → GCP Global LB transition, routing issues arose because AWS ALB still passes traffic to backends failing health checks whereas GCP LB blocks them.
Reconfigured health checks/backends after analyzing GCP NEG structure
- GKE Ingress + BackendConfig health check returned 503 when the configured requestPath had no response.
Resolved by guaranteeing the health-check endpoint response and aligning ManagedCertificate / FrontendConfig
- Game client file downloads failed after the CDN migration due to missing HTTP → HTTPS redirect.
Resolved by splitting the URL Map into HTTP / HTTPS variants (`default_url_redirect.https_redirect=true`)
- After DNS delegation to Cloud DNS, the domain was missing from Google Search Console.
Resolved by adding the issued verification TXT record to Cloud DNS to verify ownership, then re-registering

Results

- 50% cloud cost reduction (\$10,000 → \$5,000/month), 50% deployment time reduction, 95% rollback reduction
- Full GCP infrastructure Terraform IaC (TFLint/Checkov static analysis, Terratest unit testing), 100% deployment history tracking via Git commits
- Minimized downtime with resource-copy migration; final DNS cutover executed during an approximately 2-hour scheduled maintenance window aligned with game service operations
- Eliminated GCP service-account key management overhead and key-leak risk by adopting Workload Identity Federation for GitHub Actions

Tech Stack

GCP (GKE, Cloud SQL, Cloud Storage, VPC, Global LB 등), Shared VPC, Workload Identity Federation, Cloud Armor, Cloud CDN, Cloud DNS, NFS, Terraform, GitLab CI, ArgoCD, GitHub Actions, Kubernetes, Helm, Docker

Project 2: Game Data Collection Automation System

Contribution 100% 3 months (2024.01 ~ 2024.03)

Problem

Game and blockchain data was collected manually, costing analysts 2-3 hours daily, with frequent data loss caused by human error

Approach

Loaded multi-source data (MongoDB · CloudSQL · Google Analytics · Dune) into BigQuery — MongoDB via Dataflow ETL, CloudSQL via direct connection, GA/Dune via API — then triggered Daily / Monthly via Cloud Functions + Cloud Scheduler and pushed analyst-facing sheets via the Google Sheets API. A separate Cloud Function periodically deduplicated BigQuery data.

The entire infrastructure (datasets, Cloud Functions, schedulers, storage, Artifact Registry) was managed as Terraform IaC.

Results

- 95% data collection automation and 0% data gap rate by eliminating human error

Tech Stack

BigQuery, Dataflow, Cloud Functions, Cloud Scheduler, Cloud Storage, Artifact Registry, Terraform, Python, Google Sheets API, Docker

Project 3: PLG Stack Monitoring System

Contribution 100% 4 months (2024.04 ~ 2024.07)

Problem

No real-time monitoring system for Kubernetes clusters and applications, allowing only reactive responses with excessive time spent on root-cause analysis

Results

- SLI/SLO establishment and Alertmanager alarm noise optimization, proactive detection established, reduced response time, improved availability

Tech Stack

Prometheus, Loki, Grafana, Promtail, Fluent-bit, Slack API

Iorchard

2022.04 ~ 2023.03

Infra Engineer

Built PaaS (Kubernetes + OpenStack + Ceph) infrastructure solutions for clients and provided technical support, designing on-premises cloud infrastructure for financial and public-sector clients and leading operations training.

Key Projects: PaaS Solution · DP Training · Certificate Automation · Ansible Configuration Management

Results

- PaaS solution for 5 clients, Kubernetes HA 99.9% availability
- Established DP's autonomous Kubernetes operations (6-month training, avg 10 attendees · 4 hrs per session)
- 10x certificate renewal extension (1yr → 10yr), 50hrs annual ops savings
- Ansible-based server provisioning and configuration management automation eliminating configuration drift issues

Tech Stack

Kubernetes, OpenStack, Ceph, Ansible, kubeadm, Rocky Linux/CentOS

ETech System

2022.01 ~ 2022.04

System Engineer

Hardware Server and SAN storage construction. Transitioned to DevOps field for career growth.

Opensource Projects

Kubernetes CRD

[k8s-namespace-sync](#) — Automatically sync Kubernetes namespace resources across clusters
[helios-lb](#) — Custom load balancer controller for Kubernetes
[network-policy-generator](#) — Kubernetes NetworkPolicy auto-generation controller

GitHub Actions

[Compress Decompress Action](#) — Compress/decompress files in workflows
[Image Tag Updater](#) — Auto-update image tags in YAML/Helm files
[Go Git Commit Action](#) — Go-based Git commit creation and push
[Extract Commit Action](#) — Extract Git commit SHA, message, and author information
[Multi Git Mirror Action](#) — Automate mirroring across multiple Git repositories
[Kube Diff Action](#) — Compare and report Kubernetes resource changes

Ansible Galaxy

[Ansible K8s IAC Tool](#) — Automated K8s and IaC tool installation collection
[Ansible User Management](#) — Linux user and SSH key management role

DevOps Tools

[bash-pilot](#) — AI-powered Bash command assistant CLI tool
[git-bridge](#) — Git repository mirroring and sync CLI tool
[static-file-server](#) — Go-based static file server with directory UI, file preview, search, ZIP batch download

Chrome Extensions

[Dev Toolkit](#) — Developer utility collection
[QRify](#) — Instant QR code generation for URLs/text

Education

Chungnam National University - Public Administration (Transfer)	2018.03 ~ 2020.08
National Institute for Lifelong Education - Business Administration (Academic Credit Bank)	2016.09 ~ 2018.02

Certifications

Engineer Information Processing	2021.11
AWS Certified Solution Architect - Associate (Expired)	2021.11
Network Administrator Level 2	2021.10
Linux Master Level 2	2021.10
Computer Specialist Level 1	2021.04
TOEIC Speaking Test - 130/Intermediate Mid 3 (Expired)	2021.02

Study

Tech blog (somaz.blog, English) — DevOps · Kubernetes · IaC writing for global readers	2025.01 ~ 현재
Tech blog (somaz.tistory.com, Korean) — ongoing DevOps · Kubernetes · IaC learning and operations notes	2023.04 ~ 현재
T101(Terraform) Study	2023.08 ~ 2023.10
AEWS(EKS) Study	2023.04 ~ 2023.06
PKOS(Kubernetes) Study	2023.01 ~ 2023.02
Cloud Security Training	2022.02 ~ 2022.04
Security Solution Operation Specialist Training Program	2021.07 ~ 2021.12